



BIG DATA HADOOP & SPARK TRAINING

Assignment 09: Assignment on Advance Hive

Olympic Data Analysis in Hive

- Created a table "Olympic_data" with columns *name, age, country, year, to_date, sport, gold medal, silver medal, bronze medal and total of all medals*.
- Loaded the data from local filesystem to the above created Olympic_data table.
- The data set consists of the following fields.
 - ❖ Athlete: This field consists of the athlete name
 - ❖ Age: This field consists of athlete ages
 - ❖ Country: This field consists of the country names which participated in Olympics
 - ❖ Year: This field consists of the year
 - ❖ Closing Date: This field consists of the closing date of ceremony
 - ❖ Sport: Consists of the sports name
 - ❖ Gold Medals: No. of Gold medals
 - ❖ Silver Medals: No. of Silver medals
 - ❖ Bronze Medals: No. of Bronze medals
 - ❖ Total Medals: Consists of total no. of medals

Above steps are as shown below:

```
hive> create table olympic_data(name string,age int, country string,year int, to_date string,sport string, goldmedal int, silvermedal i
nt, bronzemedal int, total int) row format delimited fields terminated by ',';
OK
Time taken: 0.206 seconds      creating a table "olympicdata" ↑
hive> load data local inpath '/home/acadgild/assignments/hive/olympics_data1.csv' into table olympic_data; Loading the table with data
from local file system
FAILED: SemanticException Line 1:23 Invalid path '/home/acadgild/assignments/hive/olympics_data1.csv': No files matching path file:/h
ome/acadgild/assignments/hive/olympics_data1.csv
hive> load data local inpath '/home/acadgild/assignments/hive/olympix_data1.csv' into table olympic_data;
Loading data to table default.olympic_data
OK
```

Task 1

1. Write a Hive program to find the number of medals won by each country in swimming.

Below query will extract the total number of medals won by each country in swimming

Select country, sum(total) from Olympic_data where sport="Swimming" group by country;

```
Time taken: 1.138 seconds
hive> select country,sum(total) from olympic_data where sport="Swimming" GROUP BY country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415153417_1379b057-3897-4364-860e-bad87219ddfa
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0007, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0007/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-15 15:34:34,638 Stage-1 map = 0%, reduce = 0%
2018-04-15 15:34:49,268 Stage-1 map = 67%, reduce = 0%, Cumulative CPU 5.02 sec
2018-04-15 15:34:50,338 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.24 sec
2018-04-15 15:35:02,806 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 8.34 sec
MapReduce Total cumulative CPU time: 8 seconds 340 msec
Ended Job = job_1523776351907_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.34 sec HDFS Read: 535988 HDFS Write: 881 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 340 msec
OK
```

```
Argentina      1
Australia     163
Austria        3
Belarus        2
Brazil         8
Canada         5
China         35
Costa Rica     2
Croatia        1
Denmark        1
France        39
Germany        32
Great Britain  11
Hungary        9
Italy          16
Japan          43
Lithuania      1
Netherlands    46
Norway         2
Poland         3
Romania        6
Russia        20
Serbia         1
Slovakia       2
Slovenia       1
South Africa   11
South Korea    4
Spain          3
Sweden         9
Trinidad and Tobago 1
Tunisia        3
Ukraine        7
United States  267
Zimbabwe       7
Time taken: 46.511 seconds, Fetched: 34 row(s)
hive>
```


This is the output
for the query!

2. Write a Hive program to find the number of medals that India won year wise.

The below query will extract the total number of medal that India won every year

```
select country, year, sum(total) from olympic_data where country="India" GROUP BY country, year;
```

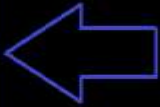
```
hive> select country, year, sum(total) from olympic_data where country="India" GROUP BY country, year;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415154541_47a47ccd-c55a-4f16-9039-7246d4dd86cd
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0008, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0008/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-15 15:45:57,796 Stage-1 map = 0%, reduce = 0%
2018-04-15 15:46:11,924 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.47 sec
2018-04-15 15:46:25,650 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 8.12 sec
MapReduce Total cumulative CPU time: 8 seconds 120 msec
Ended Job = job_1523776351907_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.12 sec HDFS Read: 536368 HDFS Write: 187 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 120 msec
OK
India 2000 1
India 2004 1
India 2008 3
India 2012 6
Time taken: 45.767 seconds, Fetched: 4 row(s)
hive>
```



3. Write a Hive Program to find the total number of medals each country won.


Select country, sum(total) from Olympic_data GROUP BY country;

```
hive> select country,sum(total) from olympic_data GROUP BY country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415154939_ffbb5b4a-1e85-4e90-bdfe-alf725302cc7
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0009, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0009/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-15 15:49:55,616 Stage-1 map = 0%, reduce = 0%
2018-04-15 15:50:07,219 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.83 sec
2018-04-15 15:50:21,018 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.76 sec
MapReduce Total cumulative CPU time: 5 seconds 760 msec
Ended Job = job_1523776351907_0009
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.76 sec HDFS Read: 535157 HDFS Write: 2742 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 760 msec
OK
Afghanistan      2
Algeria           8
Argentina        141
Armenia          10
Australia        609
Austria          91
Azerbaijan       25
Bahamas          24
Bahrain          1
Barbados         1
Belarus          97
Belgium          18
Botswana         1
```



Total number of medals each country won

```
Belgium 18
Botswana 1
Brazil 221
Bulgaria 41
Cameroon 20
Canada 370
Chile 22
China 530
Chinese Taipei 20
Colombia 13
Costa Rica 2
Croatia 81
Cuba 100
Cyprus 1
Czech Republic 81
Denmark 89
Dominican Republic 5
Ecuador 1
Egypt 8
Eritrea 1
Estonia 18
Ethiopia 29
Finland 110
France 318
Gabon 1
Georgia 23
Germany 629
Great Britain 322
Greece 59
Grenada 1
Guatemala 1
Hong Kong 3
Hungary 145
Iceland 15
India 11
Indonesia 22
Iran 24
Ireland 9
```




Output continued


```

North Korea 21
Norway 192
Panama 1
Paraguay 17
Poland 80
Portugal 9
Puerto Rico 2
Qatar 3
Romania 123
Russia 768
Saudi Arabia 6
Serbia 31
Serbia and Montenegro 38
Singapore 7
Slovakia 35
Slovenia 25
South Africa 25
South Korea 308
Spain 205
Sri Lanka 1
Sudan 1
Sweden 181
Switzerland 93
Syria 1
Tajikistan 3
Thailand 18
Togo 1
Trinidad and Tobago 19
Tunisia 4
Turkey 28
Uganda 1
Ukraine 143
United Arab Emirates 1
United States 1312
Uruguay 1
Uzbekistan 19
Venezuela 4
Vietnam 2

```

Output continued



4. Write a Hive program to find the number of gold medals each country won.

The below query will extract the total gold medal won by each country.


`select country,sum(goldmedal) from olympic_data GROUP BY country;`

```

hive> select country,sum(goldmedal) from olympic_data GROUP BY country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415155705_a2ffabd4-82eb-4f8d-b619-b177331ffe7f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0010, Tracking URL = http://localhost:8080/proxy/application_1523776351907_0010/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-15 15:57:21,092 Stage-1 map = 0%, reduce = 0%
2018-04-15 15:57:32,601 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.8 sec
2018-04-15 15:57:45,548 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.81 sec
MapReduce Total cumulative CPU time: 5 seconds 810 msec
Ended Job = job_1523776351907_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.81 sec HDFS Read: 535169 HDFS Write: 2703 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 810 msec
OK
Afghanistan 0
Algeria 2
Argentina 49
Armenia 0
Australia 163
Austria 36
Azerbaijan 6
Bahamas 11
Bahrain 0
Barbados 0
Belarus 17
Belgium 2
Botswana 0

```

Number of gold medals each country won.



```

Belgium 2
Botswana 0
Brazil 46
Bulgaria 8
Cameroon 20
Canada 168
Chile 3
China 234
Chinese Taipei 2
Colombia 2
Costa Rica 0
Croatia 35
Cuba 57
Cyprus 0
Czech Republic 14
Denmark 46
Dominican Republic 3
Ecuador 0
Egypt 1
Eritrea 0
Estonia 6
Ethiopia 13
Finland 11
France 108
Gabon 0
Georgia 6
Germany 223
Great Britain 124
Greece 12
Grenada 1
Guatemala 0
Hong Kong 0
Hungary 77
Iceland 0
India 1
Indonesia 5
Iran 10
Ireland 1

```


Output continued

```

Paraguay 0
Poland 20
Portugal 1
Puerto Rico 0
Qatar 0
Romania 57
Russia 234
Saudi Arabia 0
Serbia 1
Serbia and Montenegro 11
Singapore 0
Slovakia 10
Slovenia 5
South Africa 10
South Korea 110
Spain 19
Sri Lanka 0
Sudan 0
Sweden 57
Switzerland 21
Syria 0
Tajikistan 0
Thailand 6
Togo 0
Trinidad and Tobago 1
Tunisia 2
Turkey 9
Uganda 1
Ukraine 31
United Arab Emirates 1
United States 552
Uruguay 0
Uzbekistan 5
Venezuela 1
Vietnam 0
Zimbabwe 2
Time taken: 41.906 seconds, Fetched: 110 row(s)
hive>

```

Output continued



Task 2

Write a hive UDF that implements functionality of string concat_ws(string SEP, array<string>). This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SEP.

```
package concatws;

import org.apache.hadoop.hive.ql.exec.Description;
import org.apache.hadoop.hive.ql.exec.UDFArgumentException;
import org.apache.hadoop.hive.ql.exec.UDFArgumentLengthException;
import org.apache.hadoop.hive.ql.exec.UDFArgumentTypeException;
import org.apache.hadoop.hive.ql.metadata.HiveException;
import org.apache.hadoop.hive.ql.udf.generic.GenericUDF;
import org.apache.hadoop.hive.serde.serdeConstants;
import org.apache.hadoop.hive.serde2.objectinspector.ListObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector.Category;
import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector.PrimitiveCategory;
import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorFactory;
import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils;
import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils.PrimitiveGrouping;
import org.apache.hadoop.io.Text;

public class concatenatews extends GenericUDF {
    private transient ObjectInspector[] argumentOIs;

    @Override
    public ObjectInspector initialize(ObjectInspector[] arguments) throws UDFArgumentException {
        if (arguments.length < 2) {
            throw new UDFArgumentLengthException(
                "The function CONCAT_WS(separator,[string | array(string)]+) "
                + "needs at least two arguments.");
        }

        // check if argument is a string or an array of strings
    }
}
```



```

for (int i = 0; i < arguments.length; i++) {
    switch (arguments[i].getCategory()) {
        case LIST:
            if (isStringOrVoidType(
                ((ListObjectInspector) arguments[i]).getListElementObjectInspector())) {
                break;
            }
        case PRIMITIVE:
            if (isStringOrVoidType(arguments[i])) {
                break;
            }
        default:
            throw new UDFArgumentTypeException(i, "Argument " + (i + 1)
                + " of function CONCAT_WS must be \" + serdeConstants.STRING_TYPE_NAME
                + " or \" + serdeConstants.LIST_TYPE_NAME + "<" +
serdeConstants.STRING_TYPE_NAME
                + ">\", but \" + arguments[i].getTypeName() + \" was found.");
    }
}

argumentOIs = arguments;
return PrimitiveObjectInspectorFactory.writableStringObjectInspector;
}

protected boolean isStringOrVoidType(ObjectInspector oi) {
    if (oi.getCategory() == Category.PRIMITIVE) {
        if (PrimitiveGrouping.STRING_GROUP
            == PrimitiveObjectInspectorUtils.getPrimitiveGrouping(
                ((PrimitiveObjectInspector) oi).getPrimitiveCategory())
            || ((PrimitiveObjectInspector) oi).getPrimitiveCategory() == PrimitiveCategory.VOID)
        {
            return true;
        }
    }
    return false;
}

private final Text resultText = new Text();

@Override
public Object evaluate(DeferredObject[] arguments) throws HiveException {
    if (arguments[0].get() == null) {
        return null;
    }
    String separator = PrimitiveObjectInspectorUtils.getString(
        arguments[0].get(), (PrimitiveObjectInspector) argumentOIs[0]);

```

```

StringBuilder sb = new StringBuilder();
boolean first = true;
for (int i = 1; i < arguments.length; i++) {
    if (arguments[i].get() != null) {
        if (first) {
            first = false;
        } else {
            sb.append(separator);
        }
        if (argumentOls[i].getCategory().equals(Category.LIST)) {
            Object strArray = arguments[i].get();
            ListObjectInspector strArrayOI = (ListObjectInspector) argumentOls[i];
            boolean strArrayFirst = true;
            for (int j = 0; j < strArrayOI.getListLength(strArray); j++) {
                if (strArrayFirst) {
                    strArrayFirst = false;
                } else {
                    sb.append(separator);
                }
                sb.append(strArrayOI.getListElement(strArray, j));
            }
        } else {
            sb.append(PrimitiveObjectInspectorUtils.getString(
                arguments[i].get(), (PrimitiveObjectInspector) argumentOls[i]));
        }
    }
}

resultText.set(sb.toString());
return resultText;
}

@Override
public String getDisplayString(String[] children) {
    assert (children.length >= 2);
    return getStandardDisplayString("concat_ws", children);
}
}

```

- The above program is written in eclipse and exported as jar named **“newHiveUDFtask1.jar”**
- This jar needs to be added to Hive, this can be done by using the command in hive “*add jar <local path of jar file saved>*”
- We have to create a temporary function to use this function only for the current instance of hive. This can be done using the command “*create temporary function <class_name> <'package_name.class_name'>*”
- We have a table in hive “people” which has two columns “name, friend’s name” this table represents name of a person and his/her friend’s name.

```
hive> select * from people;
OK
bob      sara
bob      john
bob      ted
john     sara
ted      bob
ted      sara
Time taken: 5.868 seconds, Fetched: 6 row(s)
hive>
hive> █
```

- We will use the above UDF concatenatews to concatenate the above data with “,” separator.

Above steps are as shown below:

```
hive> create table people(name string, friends string) row format delimited fields terminated by '\t';
OK
Time taken: 0.126 seconds
hive> load data local inpath '/home/acadgild/assignments/hive/hiveudf.txt' into table people;
Loading data to table default.people
OK
Time taken: 0.85 seconds
hive> add jar /home/acadgild/assignments/hive/newHiveUDFtask1.jar;
Added [/home/acadgild/assignments/hive/newHiveUDFtask1.jar] to class path
Added resources: [/home/acadgild/assignments/hive/newHiveUDFtask1.jar]
hive> create temporary function concatenatews as 'concatws.concatenatews';
OK
Time taken: 0.054 seconds
hive> select concatenatews(',',name,friends) from people;
OK
bob,sara
bob,john
bob,ted
john,sara
ted,bob
ted,sara
Time taken: 0.237 seconds, Fetched: 6 row(s)
hive> █
```

Created the table and loaded the data into the table

Adding the jar to Hive and creating a temporary function

concatenated output

Using the temporary function created to concatenate the data with “,” separator

Task 3

Link: <https://acadgild.com/blog/transactions-in-hive/>

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

The different row-level transactions available in Hive are as follows:

1. Insert
2. Delete
3. Update

Row-level Transactions Available in Hive:

Before creating a Hive table that supports transactions, the transaction features present in Hive needs to be turned on, as by default they are turned off.

The below properties needs to be set appropriately in *hive shell*, order-wise to work with transactions in Hive:

```
hive>set hive.support.concurrency = true;
```

```
hive>set hive.enforce.bucketing = true;
```

```
hive>set hive.exec.dynamic.partition.mode = nonstrict;
```

```
hive>set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
```

```
hive>set hive.compactor.initiator.on = true;
```

```
hive>set hive.compactor.worker.threads = a positive number on at least one instance of the Thrift metastore service;
```

```
hive> set hive.support.concurrency = true;
hive> set hive.enforce.bucketing = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive> set hive.compactor.initiator.on = true;
```

Creating a Table That Supports Hive Transactions

```
CREATE TABLE college(clg_id int,clg_name string,clg_loc string) clustered by (clg_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
```

The above syntax will create a table with name 'college' and the columns present in the table are 'clg_id, clg_name, clg_loc'. We are *bucketing* the table by 'clg_id' and the table format is 'orc', also we are enabling the transactions in the table by specifying it inside the *TBLPROPERTIES* as 'transactional'='true'

```
hive> CREATE TABLE college(clg_id int,clg_name string,clg_loc string) clustered by (clg_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
OK
Time taken: 0.848 seconds
```


Inserting Data into a Hive Table

```
insert into table college
```

```
values(1,'SSMRV','Jayanagar'),(2,'RVENG','Kengeri'),(3,'PESIT','MysoreRd'),(4,'CMRIT','Kundenalli Gate'),(5,'AMC','Bommasandra');
```


The above command is used to insert row wise data into the Hive table. Here, each row is separated by '()' brackets. The contents of the table can be viewed using the command *select * from college*

```
hive> insert into table college values(1,'SSMRV','Jayanagar'),(2,'RVENG','Kengeri'),(3,'PESIT','MysoreRd'),(4,'CMRIT','Kundenalli Gate'),(5,'AMC','Bommasandra');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415181908_1006ecel-c0d5-474f-896e-e12e42499653
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0011, Tracking URL = http://localhost:8080/proxy/application_1523776351907_0011/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0011
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-04-15 18:19:25,133 Stage-1 map = 0%, reduce = 0%
2018-04-15 18:19:37,925 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 5.03 sec
2018-04-15 18:20:17,751 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 7.28 sec
2018-04-15 18:20:28,032 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 8.97 sec
2018-04-15 18:20:31,014 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 10.65 sec
2018-04-15 18:20:32,412 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 12.3 sec
2018-04-15 18:20:33,824 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 16.38 sec
2018-04-15 18:20:46,386 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 21.13 sec
2018-04-15 18:20:52,635 Stage-1 map = 100%, reduce = 87%, Cumulative CPU 25.06 sec
2018-04-15 18:20:55,420 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 30.38 sec
2018-04-15 18:20:57,823 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 34.79 sec
MapReduce Total cumulative CPU time: 34 seconds 790 msec
Ended Job = job_1523776351907_0011
Loading data to table default.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 34.79 sec HDFS Read: 27078 HDFS Write: 4098 SUCCESS
Total MapReduce CPU Time Spent: 34 seconds 790 msec
OK
Time taken: 111.574 seconds
hive> select * from college;
OK
5      AMC      Bommasandra
1      SSMRV     Jayanagar
2      RVENG     Kengeri
3      PESIT     MysoreRd
4      CMRIT     Kundenalli Gate
Time taken: 0.332 seconds, Fetched: 5 row(s)
hive>
```



we will re-insert the same data again, it will be appended to the previous data as shown below:

```
hive> insert into table college values(1,'SSMRV','Jayanagar'),(2,'RVENG','Kengeri'),(3,'PESIT','MysoreRd'),(4,'CMRIT','Kundenalli Gate'),(5,'AMC','Bommasandra');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415182449_707aeee6-9951-4629-b810-64ec280c8d12
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0012, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0012/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0012
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-04-15 18:25:10,113 Stage-1 map = 0%, reduce = 0%
2018-04-15 18:25:22,858 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.59 sec
2018-04-15 18:26:01,212 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 4.72 sec
2018-04-15 18:26:07,212 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 6.93 sec
2018-04-15 18:26:08,636 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 8.04 sec
2018-04-15 18:26:10,618 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 9.14 sec
2018-04-15 18:26:14,410 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 12.12 sec
2018-04-15 18:26:25,199 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 15.72 sec
2018-04-15 18:26:29,619 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 19.88 sec
2018-04-15 18:26:33,739 Stage-1 map = 100%, reduce = 87%, Cumulative CPU 24.28 sec
2018-04-15 18:26:34,981 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 32.35 sec
MapReduce Total cumulative CPU time: 32 seconds 550 msec
Ended Job = job_1523776351907_0012
Loading data to table default.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 32.55 sec HDFS Read: 26863 HDFS Write: 4098 SUCCESS
Total MapReduce CPU Time Spent: 32 seconds 550 msec
OK
Time taken: 108.618 seconds
hive> select * from college;
OK
5      AMC      Bommasandra
5      AMC      Bommasandra
1      SSMRV     Jayanagar
1      SSMRV     Jayanagar
2      RVENG     Kengeri
2      RVENG     Kengeri
3      PESIT     MysoreRd
3      PESIT     MysoreRd
4      CMRIT     Kundenalli Gate
4      CMRIT     Kundenalli Gate
Time taken: 0.323 seconds, Fetched: 10 row(s)
hive>
```



Values have got appended to the previous values after performing the same insert operation again

Updating the Data in Hive Table

UPDATE college set clg_id = 6 where clg_id = 3; // not supported because of bucketing

we can see that we have received an error message. This means that the Update command is not supported on the columns that are bucketed. we have bucketed the 'clg_id' column and performing the Update operation on the same column, so we got the error.

But we can perform **update operation on Non bucketed column**

UPDATE college set clg_name = 'IIT' where clg_id = 2;

The updated data can be checked using the command *select * from college.*

```
hive> UPDATE college set clg_id = 6 where clg_id = 3;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is not supported. Column clg_id.
hive> UPDATE college set clg_name = 'IIT' where clg_id = 2;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415183142_3d9e6769-8dc7-4e0d-aaa6-6005fafa2f64
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0013, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0013/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0013
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-04-15 18:31:59,504 Stage-1 map = 0%, reduce = 0%
2018-04-15 18:32:59,012 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 13.91 sec
2018-04-15 18:33:07,596 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 16.39 sec
2018-04-15 18:33:08,889 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 19.03 sec
2018-04-15 18:33:10,240 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 26.87 sec
2018-04-15 18:33:53,773 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 27.9 sec
2018-04-15 18:33:56,863 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 33.19 sec
2018-04-15 18:33:58,166 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 34.74 sec
2018-04-15 18:34:06,955 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 36.61 sec
2018-04-15 18:34:08,122 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 42.05 sec
2018-04-15 18:34:09,192 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 44.24 sec
MapReduce Total cumulative CPU time: 44 seconds 240 msec
Ended Job = job_1523776351907_0013
Loading data to table default.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 44.24 sec HDFS Read: 56321 HDFS Write: 996 SUCCESS
Total MapReduce CPU Time Spent: 44 seconds 240 msec
OK
Time taken: 148.512 seconds
hive> select * from college;
OK
5      AMC      Bommasandra
5      AMC      Bommasandra
1      SSMRV     Jayanagar
1      SSMRV     Jayanagar
2      IIT      Kengeri
2      IIT      Kengeri
3      PESIT     MysoreRd
3      PESIT     MysoreRd
4      CMRIT     Kundenalli Gate
4      CMRIT     Kundenalli Gate
Time taken: 0.248 seconds, fetched: 10 row(s)
hive>
```

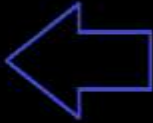
Updated data

Deleting a Row from Hive Table:

delete from college where clg_id=5;

```
hive> delete from college where clg_id=5;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180415183918_7b48f6ac-6b03-4c8b-b807-1131f7236550
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1523776351907_0014, Tracking URL = http://localhost:8088/proxy/application_1523776351907_0014/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1523776351907_0014
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-04-15 18:39:33,873 Stage-1 map = 0%, reduce = 0%
2018-04-15 18:40:33,686 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 15.57 sec
2018-04-15 18:40:36,737 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 17.66 sec
2018-04-15 18:40:41,173 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 19.81 sec
2018-04-15 18:40:42,738 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 21.96 sec
2018-04-15 18:40:44,163 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 23.96 sec
2018-04-15 18:41:23,706 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 25.6 sec
2018-04-15 18:41:25,245 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 28.73 sec
2018-04-15 18:41:26,689 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 29.76 sec
2018-04-15 18:41:29,418 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 32.15 sec
2018-04-15 18:41:35,145 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 35.36 sec
2018-04-15 18:41:36,545 Stage-1 map = 100%, reduce = 87%, Cumulative CPU 37.93 sec
2018-04-15 18:41:37,764 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 41.12 sec
MapReduce Total cumulative CPU time: 41 seconds 120 msec
Ended Job = job_1523776351907_0014
Loading data to table default.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 41.12 sec HDFS Read: 54468 HDFS Write: 755 SUCCESS
Total MapReduce CPU Time Spent: 41 seconds 120 msec
OK
Time taken: 141.345 seconds
hive> select * from college;
OK
1      SSMRV  Jayanagar
1      SSMRV  Jayanagar
2      IIT    Kengeri
2      IIT    Kengeri
3      PESIT  MysoreRd
3      PESIT  MysoreRd
4      CMRIT  Kundenalli Gate
4      CMRIT  Kundenalli Gate
Time taken: 0.276 seconds, Fetched: 8 row(s)
hive>
```

Values of collegeid 5 is deleted



We have now successfully deleted a row from the Hive table. This can be checked using the command **select * from college**, as shown above. We can see that there is no row with **clg_id = 5**.